

Deep Kalman Filters

Chizhov Nikita Ekaterina Ivanova Vadim Liventsev

Skolkovo Institute of Science and Technology

October 26, 2018

Kalman Filter




Figure: Model scheme

Goal. Fit generative model to a sequence of observations and actions.

- Observations $\vec{x} = (x_1, \dots, x_T)$, $x_t \in \mathbb{R}^d$, noisy, non-linear function of latent space
- Actions $\vec{u} = (u_1, \dots, u_{T-1})$, $u_t \in \mathbb{R}^c$
- Latent space $\vec{z} = (z_1, \dots, z_T)$, $z_t \in \mathbb{R}^s$

$$\begin{aligned}z_1 &\sim \mathcal{N}(\mu_0; \Sigma_0) \\z_t &\sim \mathcal{N}(G_\alpha(z_{t-1}, u_{t-1}, \delta_t), S_\beta(z_{t-1}, u_{t-1}, \delta_t)) \\x_t &\sim \Pi(F_\kappa(z_t)).\end{aligned}$$

$G_\alpha, S_\beta, F_\kappa$ are nonlinear functions of the previous state, parametrized by deep neural networks.

Distribution of the observations x_t are Multinomial.

Parameters of the generative model $\theta = \alpha, \beta, \kappa$.

Maximizing ELBO

We would like to fit generative model w.r.t. parameters θ .

$$\max_{\theta} \log_{\theta} p(x_1, \dots, x_T | u_1, \dots, u_{T-1})$$


Using variational inference we reduce the problem to optimization of lower bound.

$$q_{\phi}(\vec{z} | \vec{x}, \vec{u}) = \prod_{t=1}^T q(z_t | z_{t-1}, x_1, \dots, x_T, \vec{u}) \sim \text{prior}$$

ELBO of conditional log likelihood:

$$\begin{aligned} \log p_{\theta}(\vec{x} | \vec{u}) &\geq \mathcal{L}(x; (\theta, \phi)) = \\ &\sum_{t=1}^T \mathbb{E}_{q_{\phi}(z_t | \vec{x}, \vec{u})} [\log p_{\theta}(x_t | z_t)] - \text{KL}(q_{\phi}(z_1 | \vec{x}, \vec{u}) || p_0(z_1)) \\ &- \sum_{t=2}^T \mathbb{E}_{q_{\phi}(z_{t-1} | \vec{x}, \vec{u})} [\text{KL}(q_{\phi}(z_t | z_{t-1}, \vec{x}, \vec{u}) || p_0(z_t | z_{t-1}, u_{t-1}))]. \end{aligned}$$

Algorithm



(a) Deep Kalman Filter

Algorithm 1 Learning Deep Kalman Filters

```
while notConverged() do  
   $\vec{x} \leftarrow \text{sampleMiniBatch}()$   
  Perform inference and estimate likelihood:  
  1.  $\hat{z} \sim q_\phi(\vec{z} | \vec{x}, \vec{u})$   
  2.  $\hat{x} \sim p_\theta(\vec{x} | \hat{z})$   
  3. Compute  $\nabla_\theta \mathcal{L}$  and  $\nabla_\phi \mathcal{L}$  (Differentiating  
  4. Update  $\theta, \phi$  using ADAM  
end while
```

(b) Algorithm for learning DKF

State transitions $z_t \rightarrow z_{t+1}$ are parametrized by a 2-layer stacked LSTM
Mean of normal distribution $p(\vec{x}|\vec{z})$ is parametrized by a multilayer perceptron

Mean of normal distribution $q(\vec{z}|\vec{x})$ is one of the following

- *q-INDEP*: $q(z_t|x_t, u_t)$ parameterized by an MLP
- *q-LR*: $q(z_t|x_{t1}, x_t, x_{t+1}, u_{t1}, u_t, u_{t+1})$ parameterized by an MLP
- *q-RNN*: $q(z_t|x_1, \dots, x_t, u_1, \dots, u_t)$ parameterized by a RNN
- *q-BRNN*: $q(z_t|x_1, \dots, x_T, u_1, \dots, u_T)$ parameterized by a bi-directional RNN

Dataset

For our project we used **Healing MNIST**: a synthetic dataset is obtained by producing sequence of rotated images.

- Actions \vec{u} are the rotations
- Observations \vec{x} are the rotated images.

We have *60000 of training sequences*, each sequence has *5 rotated images*.

To one sequence of three consecutive squares is superimposed with the top-left corner of the images.




Figure: Sequence example.

Rotation encoding

In the reproduced paper authors don't specify how they represent a rotation with \vec{u} ¹

We used an encoding scheme inspired by *Attention is all you need*

$$s_i = \sin\left(\frac{\pi}{2} \frac{\theta}{\theta_{max}} i\right), \quad i \in [1, 25] \quad (1)$$

$$c_i = \cos\left(\frac{\pi}{2} \frac{\theta}{\theta_{max}} i\right), \quad i \in [1, 25] \quad (2)$$

$$\vec{u} = (\vec{s}, \vec{c}) \quad (3)$$

¹Maybe that's the secret sauce that makes their model work?

We implemented

- ① Healing MNIST
- ② Kalman filter architecture proposed by the original authors
- ③ Stacked LSTM for the transitional model and 2 options (q -INDEP and q -RNN) for the inference model
- ④ Simple Variational Autoencoder
- ⑤ Simple Autoencoder

The original results *did not reproduce*. Simple autoencoder works best on Healing MNIST

Reconstructions




Figure: Simple autoencoder




Figure: Deep Kalman filters, variational autoencoder

References



Krishnan, Rahul G., Uri Shalit, and David Sontag. "Deep kalman filters." *arXiv preprint arXiv:1511.05121* (2015)



Krishnan, Rahul G., Uri Shalit, and David Sontag. "Structured Inference Networks for Nonlinear State Space Models." *AAAI*. 2017.