
DEEP KALMAN FILTERS

A PREPRINT

Nikita Chizhov

Skolkovo Institute of Science and Technology
Nikita.Chizhov@skoltech.ru

Ekaterina Ivanova

Skolkovo Institute of Science and Technology
Ekaterina.Ivanova@skoltech.ru

Vadim Liventsev

Skolkovo Institute of Science and Technology
Vadim.Liventsev@skoltech.ru

October 29, 2018

Keywords Kalman Filter · Deep learning · Bayesian methods

1 Introduction

In Deep Kalman Filters, Krishnan et al. attempt to apply a modified version of Kalman Filters to determine what the best course of treatment for a patient might be, even for patients with only an incomplete medical history. Kalman filters are generative probabilistic graphical models for time series data. Despite the broad success of Kalman Filters, they are limited in their ability to model complex sequential data due to their assumptions of linear relationships between successive latent states and between latent states and observed states. Extension of classical Kalman Filters by replacing linear transformations with nonlinear transformations parameterized by neural networks, is proposed in “Deep Kalman Filters” [1]. The technique was applied to perturbed MNIST data, which the authors refer to as a "Healing MNIST database".

2 Background

The classical Kalman Filter aims to model latent states $\vec{z} = (z_1, \dots, z_T)$, $x_t \in \mathbb{R}^s$, using a sequence of observations $\vec{x} = (x_1, \dots, x_T)$, $x_t \in \mathbb{R}^d$ and actions $\vec{u} = (u_1, \dots, u_{T-1})$, $u_t \in \mathbb{R}^c$, as follows:

$$z_t = \mathcal{N}(G_t z_{t-1} + B_t u_{t-1} + \epsilon_t \text{ (action-transition)}), \quad x_t = F_t z_t + \eta_t \text{ (observation)} \quad (1)$$

3 Model

The authors propose to change all the linear functions the Kalman filter modeling to be neural networks. Suppose that we have unobserved latent variables $\vec{z} = (z_1, \dots, z_T)$ and corresponding observations $\vec{x} = (x_1, \dots, x_T)$, $x_t \in \mathbb{R}^d$ and actions $\vec{u} = (u_1, \dots, u_{T-1})$, $u_t \in \mathbb{R}^c$, $x_t \in \mathbb{R}^s$. We assume the observations are a noisy, non-linear function of this latent state. We also assume that we can observe actions, which affect the latent state in a non-linear manner.

The generative model for the deep Kalman filter is then given by:

$$\begin{aligned} z_1 &\sim \mathcal{N}(\mu_0; \Sigma_0) \\ z_t &\sim \mathcal{N}(G_\alpha(z_{t-1}, u_{t-1}, \delta_t), S_\beta(z_{t-1}, u_{t-1}, \delta_t)) \\ x_t &\sim \Pi(F_\kappa(z_t)). \end{aligned} \quad (2)$$

The observations x_t are distributed according to a distribution Π - a Bernoulli distribution if the data is binary, whose parameters are a function of the corresponding latent state z_t . The functions $G_\alpha; S_\beta; F_\kappa$ are assumed to be parameterized by deep neural networks. We set $\mu_0 = 0, \Sigma_0 = I_d$.

3.1 ELBO maximization

Our training step involves optimizing our choices of F, G , and B such that they maximize the conditional log-likelihood $\max_\theta \log p_\theta(x)$ for a generative model. This can be done tractably, if not analytically, for linear transformations. However, non-linear transformations render the posterior distribution $p_\theta(\vec{z}|\vec{x}, \vec{z})$ intractable to compute. In order to overcome the intractability of posterior inference, we make use of variational autoencoders. The key concept is that an approximation q_ϕ of the posterior is trained using a recognition network.

We want to maximize the lower bound $\mathcal{L}(x; (\phi, \theta))$ in the following equation using stochastic backpropagation:

$$\begin{aligned} \log p_\theta(\vec{x}|\vec{z}) &\geq \mathcal{L}(x; (\phi, \theta)) = \\ &\sum_{t=1}^T \mathbb{E}_{q_\phi(z_t|\vec{x}, \vec{z})} [p_\theta(x_t|z_t)] - KL(q_\phi(z_1|\vec{x}, \vec{z})||p_0(z_1)) - \\ &\sum_{t=2}^T \mathbb{E}_{q_\phi(z_{t-1}|\vec{x}, \vec{z})} KL(q_\phi(z_t|z_{t-1}, \vec{x}, \vec{z})||p_0((z_t|z_{t-1}, u_{t-1}))) \end{aligned}$$

ELBO is differentiable in the parameters of the model $(\phi; \theta)$, and we can apply backpropagation for updating θ , and stochastic backpropagation for estimating the gradient w.r.t. ϕ of the expectation terms w.r.t. $q_\phi(z_t)$. Since we can use Monte Carlo estimates of the gradients of $\mathbb{E}_{q_\phi(z_t|\vec{x}, \vec{z})} [p_\theta(x_t|z_t)]$ and $KL(q_\phi(z_1|\vec{x}, \vec{z})||p_0(z_1))$ w.r.t. $q_\phi(z_t)$, we can use stochastic back-propagation in a variational autoencoder whose objective is to maximize the lower bound using an inference-update algorithm (fig.1)..

Algorithm 1 Learning Deep Kalman Filters

```

while not Converged() do
   $\vec{x} \leftarrow \text{sampleMiniBatch}()$ 
  Perform inference and estimate likelihood:
  1.  $\hat{z} \sim q_\phi(\vec{z}|\vec{x}, \vec{u})$ 
  2.  $\hat{x} \sim p_\theta(\vec{x}|\hat{z})$ 
  3. Compute  $\nabla_\theta \mathcal{L}$  and  $\nabla_\phi \mathcal{L}$  (Differentiating
  4. Update  $\theta, \phi$  using ADAM
end while

```

Figure 1: Algorithm for learning DKF

In this work we use several types of recognition network, which were proposed by the authors. They consider four variational models: two multi-layer perceptron and two neural network models.

- *q-INDEP*: $q(z_t|x_t, u_t)$ parameterized by an MLP
- *q-LR*: $q(z_t|x_{t-1}, x_t, x_{t+1}, u_{t-1}, u_t, u_{t+1})$ parameterized by an MLP
- *q-RNN*: $q(z_t|x_1, \dots, x_t, u_1, \dots, u_t)$ parameterized by a RNN
- *q-BRNN*: $q(z_t|x_1, \dots, x_T, u_1, \dots, u_T)$ parameterized by a bi-directional RNN

4 Experiments

4.1 Dataset

In our project we tried to reproduce the original dataset used in the paper. Authors called synthetic dataset ‘‘Healing MNIST’’. To create this dataset, we used images from the original MNIST dataset and subjected them to step by step

transformations (rotations), as well as bit flipping with 15%, to obtain a sequence of transformations of the same image. We also added superimposed an image over the top left corner of three consecutive transformations as it was proposed in original paper. We made experiments on several versions of created dataset. The first one consists of sequences of transformations of images of all digits, and the smaller one, which contains sequences of transformations of images of only one digit 1. There were 60000 and 6000 sequences respectively in each dataset.




Figure 2: Healing MNIST

4.2 Model variations

Due to a number of details being omitted in the original paper we experimented with different options for filling these blanks. We tried:

Representing G_α as:

- MLP with 1, 2, 3, 4, 5 layers with ReLU nonlinearity
- 1 and 2 layered GRU (with z_{t-1} as hidden state and u_{t-1} as input)
- 1 and 2 layered LSTM (with z_{t-1} as a concatenation of hidden state and cell state and u_{t-1} as input). We tried hidden layer sizes 784 (the size of an input image), 500, 200, and 80

Representing F_κ as MLP with 1, 2, 3, 4, 5 with ReLU nonlinearity and sigmoid on top of the last layer

We implemented 2 variational models (see section 3.1) $q(z|x, u)$:

1. $q_{\text{INDEP}}(z_t|x_t, u_t)$, parametrized with a 5-layer fully connected neural network with ReLU or Tanh (we tried both) $Q(x, u)$
2. $q_{\text{RNN}}(z_t|x, u)$, parametrized with a Recurrent Neural Network $Q(x, u)$

The best reconstruction of MNIST via any of those models is the following:




Figure 3: Best reconstruction with Deep Kalman Filters

5 Training procedure

In the *forward pass* we

1. Sample a batch of images x and actions u from the dataset
2. Calculate the mean of $q(z|x, u)$ using $Q(x, u)$
3. Draw a sample from the standard normal distribution and multiply by mean and variance of $q(z|x, u)$ to obtain sample z (reparametrization trick)
4. Apply $G_\alpha(z_t)$ to z_1, z_2, \dots, z_{n-1} to obtain $\hat{z}_2, \hat{z}_3, \dots, \hat{z}_n$
5. Calculate ELBO using the formula from section 3.1

In the *backward pass* we use Pytorch to backpropagate ELBO through our entire model and make an update to the model parameters using Adam optimizer with learning rate 0.001.

6 Baseline model

For completeness sake, in addition to various Deep Kalman Filter models, we implemented a model that works: a simple (non-variational) autoencoder. Its encoder is comprised of 4 fully connected layers with ReLU activations (size 784-400-200-100-50) whilst the decoder architecture is a precise reverse: 50-100-200-400-784 with ReLU activations. We train this model to encode every image separately, ignoring the links between related images in a sequence. See figure 4 for image reconstruction examples.

7 Team members contribution

- **Nikita Chizhov**: create experiments template, conducting experiments
- **Ekaterin Ivanova**: conducting experiments, presentation preparation, preparation of report
- **Vadim Liventsev**: creation of the dataset, conducting experiments, presentation preparation

References

- [1] Rahul G Krishnan, Uri Shalit, and David Sontag. Deep kalman filters. *arXiv preprint arXiv:1511.05121*, 2015.




Figure 4: Reconstruction with Simple Autoencoder